

Lab 3 - Data Intensive Computing

Stream Processing with Flink

Windows and Watermarks

In an ideal stream, every time the end of a window is hit, one can assume that all the records belonging to that window have arrived and hence the processing of that window can be started immediately. However, due to issues such as network latency and processing power, it can happen that records arrive past the closing time of a window. To solve this issue, the sources of these events can emit watermarks: when a watermark is received we can assume that all the records before the watermark have been observed. With this indication, it is enough to wait for the right watermark before processing the contents of a window.

Preprocessing

With this preprocessing, the latitude and longitude of each ride are converted into cell ids, then only the rides that started or ended in an terminal of the airport are kept.

```
private val extractLatLon = (ride: TaxiRide) =>
  if (ride.isStart)
    (ride.startLon, ride.startLat)
  else
    (ride.endLon, ride.endLat)
```

```
private val extractTerminal = (lon: Float, lat: Float) =>
  TerminalUtils.gridToTerminal(GeoUtils.mapToGridCell(lon, lat))
```

```
val maxDelay = 60 // events are out of order by max 60 seconds
val speed = 600 // events of 10 minutes are served in 1 second
```

```
val env = StreamExecutionEnvironment.getExecutionEnvironment
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime)
val rides = env.addSource(new TaxiRideSource("nycTaxiRides.gz", maxDelay,
  speed))
```

```
val jfkRides = rides
  .map(extractLatLon)
  .map(extractTerminal.tupled)
  .filter(_ != Terminal_404)
```

Problem 1 - Number of taxis in each terminal by hour

The stream of ride events is grouped by terminal and windowed with a tumbling window of one hour. Then, every terminal-window pair is processed to output a tuple with the terminal, the ride events count and the start of the window in hours.

```
val windowedCountsByTerminal = jfkRides
  .keyBy(terminal => terminal)
  .window(TumblingEventTimeWindows.of(Time.hours(1)))
  .apply((terminal: Terminal, window, records, out: Collector[(Terminal, Int, Int)]) => {
    val cal: Calendar = Calendar.getInstance()
    cal.setTimeZone(TimeZone.getTimeZone("America/New_York"))
    cal.setTimeInMillis(window.getStart)
    val hour = cal.get(Calendar.HOUR_OF_DAY)
    out.collect((terminal, records.size, hour))
  })
windowedCountsByTerminal.print()
```

Problem 2 - Busiest terminal by hour

In order to find the busiest terminal for every hour, the stream of rides is windowed without grouping on the terminal id. Then, for every window a count of rides per terminal is computed. Finally, only the terminal with the highest rides count is emitted.

```
val windowedBusiestTerminal = jfkRides
  .timeWindowAll(Time.hours(1))
  .apply((window, records, out: Collector[(Terminal, Int, Int)]) => {
    val cal: Calendar = Calendar.getInstance()
    cal.setTimeZone(TimeZone.getTimeZone("America/New_York"))
    cal.setTimeInMillis(window.getStart)
    val hour = cal.get(Calendar.HOUR_OF_DAY)
    val counts = records
      .foldLeft(HashMap.empty[Terminal, Int].withDefaultValue(0))((acc, terminal) => {
        acc(terminal) += 1
        acc
      })
    val busiestTerminal = counts.maxBy(_._2)
    out.collect((busiestTerminal._1, busiestTerminal._2, hour))
  })
windowedBusiestTerminal.print()
```

JfkTerminals.scala

```

1  package lab3.problems
2
3  import java.util.{Calendar, TimeZone}
4
5  import com.dataartisans.flinktraining.exercises.datastream_java.datatypes.TaxiRide
6  import com.dataartisans.flinktraining.exercises.datastream_java.sources.TaxiRideSource
7  import com.dataartisans.flinktraining.exercises.datastream_java.utils.GeoUtils
8  import org.apache.flink.streaming.api.TimeCharacteristic
9  import org.apache.flink.streaming.api.scala._
10 import org.apache.flink.streaming.api.windowing.assigners.TumblingEventTimeWindows
11 import org.apache.flink.streaming.api.windowing.time.Time
12 import org.apache.flink.util.Collector
13
14 import scala.collection.mutable.HashMap
15
16 object JfkTerminals {
17   private val extractLatLon = (ride: TaxiRide) =>
18     if (ride.isStart)
19       (ride.startLon, ride.startLat)
20     else
21       (ride.endLon, ride.endLat)
22
23   private val extractTerminal = (lon: Float, lat: Float) =>
24     TerminalUtils.gridToTerminal(GeoUtils.mapToGridCell(lon, lat))
25
26   def main(args: Array[String]) {
27     val maxDelay = 60 // events are out of order by max 60 seconds
28     val speed = 600 // events of 10 minutes are served in 1 second
29
30     val env = StreamExecutionEnvironment.getExecutionEnvironment
31     env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime)
32     val rides = env.addSource(new TaxiRideSource("nycTaxiRides.gz", maxDelay, speed))
33
34     val jfkRides = rides
35       .map(extractLatLon)
36       .map(extractTerminal.tupled)
37       .filter(_ != Terminal_404)
38
39     val windowedCountsByTerminal = jfkRides
40       .keyBy(terminal => terminal)
41       .window(TumblingEventTimeWindows.of(Time.hours(1)))
42       .apply((terminal: Terminal, window, records, out: Collector[(Terminal, Int, Int)]) => {
43         val cal: Calendar = Calendar.getInstance()
44         cal.setTimeZone(TimeZone.getTimeZone("America/New_York"))
45         cal.setTimeInMillis(window.getStart)
46         val hour = cal.get(Calendar.HOUR_OF_DAY)
47         out.collect((terminal, records.size, hour))
48       })
49     windowedCountsByTerminal.print()
50
51     val windowedBusiestTerminal = jfkRides
52       .timeWindowAll(Time.hours(1))
53       .apply((window, records, out: Collector[(Terminal, Int, Int)]) => {
54         val cal: Calendar = Calendar.getInstance()
55         cal.setTimeZone(TimeZone.getTimeZone("America/New_York"))
56         cal.setTimeInMillis(window.getStart)
57         val hour = cal.get(Calendar.HOUR_OF_DAY)
58
59         val counts = records
60           .foldLeft(HashMap.empty[Terminal, Int].withDefaultValue(0))((acc, terminal) => {
61             acc(terminal) += 1
62             acc
63           })
64         val busiestTerminal = counts.maxBy(_._2)
65         out.collect((busiestTerminal._1, busiestTerminal._2, hour))
66       })
67     windowedBusiestTerminal.print()
68
69     // execute program
70     env.execute("Airport Terminals")
71   }
72 }

```

TerminalUtils.scala

```
1  package lab3.problems
2
3  sealed trait Terminal{def grid: Int}
4  case object Terminal_1 extends Terminal {val grid = 71436}
5  case object Terminal_2 extends Terminal {val grid = 71688}
6  case object Terminal_3 extends Terminal {val grid = 71191}
7  case object Terminal_4 extends Terminal {val grid = 70945}
8  case object Terminal_5 extends Terminal {val grid = 70190}
9  case object Terminal_6 extends Terminal {val grid = 70686}
10 case object Terminal_404 extends Terminal {val grid = -1}
11
12 object TerminalUtils {
13     val terminals: Set[Terminal] = Set(Terminal_1, Terminal_2, Terminal_3, Terminal_4, Terminal_5, Terminal_6)
14
15     def gridToTerminal(gridCell: Int): Terminal = {
16         terminals.find(t => t.grid == gridCell) match {
17             case Some(terminal) => terminal;
18             case None => Terminal_404;
19         }
20     }
21 }
22
23
```